

EEE413 Data Communication and Communication Networks

Dr Kyeong Soo (Joseph) Kim

Department of Electrical and Electronic Engineering
Xi'an Jiaotong-Liverpool University

16 September 2019

Outline

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System

- Other Markov Systems

Network Simulation

- Ergodicity

- Time-Stepped vs. Discrete-Event Simulation

- Discrete-Event Simulation

- SimPy: Python-Based Discrete-Event Simulation Framework

Next ...

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System

- Other Markov Systems

Network Simulation

- Ergodicity

- Time-Stepped vs. Discrete-Event Simulation

- Discrete-Event Simulation

- SimPy: Python-Based Discrete-Event Simulation

- Framework

Basic Queueing Theory

Example: Queueing in An Airport



Example: Queueing in An Airport



Next ...

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System

- Other Markov Systems

Network Simulation

- Ergodicity

- Time-Stepped vs. Discrete-Event Simulation

- Discrete-Event Simulation

- SimPy: Python-Based Discrete-Event Simulation

- Framework


Little's Theorem: Definitions - State Variables

- ▶ $N(t)$: Number of customers in the system at time t .
- ▶ $\alpha(t)$: Number of customers arrived during the interval $[0, t]$.
- ▶ T_i : Time spent in the system by the i th customer.

Little's Theorem: Definitions - Steady States

- ▶ Steady-state number of customers: $N \triangleq \lim_{t \rightarrow \infty} N_t$,
where $N_t = \frac{1}{t} \int_0^t N(\tau) d\tau$.
- ▶ Steady-state arrival rate: $\lambda \triangleq \lim_{t \rightarrow \infty} \lambda_t$,
where $\lambda_t = \frac{\alpha(t)}{t}$.
- ▶ Steady-state customer delay: $T \triangleq \lim_{t \rightarrow \infty} T_t$,
where $T_t = \frac{\sum_{i=1}^{\alpha(t)} T_i}{\alpha(t)}$.

Little's Theorem: Graphical Proof I



Little's Theorem: Graphical Proof II

- ▶ Because the shaded area (up to time t) can be expressed as either $\int_0^t N(\tau)d\tau$ (in the right figure) or $\sum_{i=1}^{\alpha(t)} T_i$ (in the left figure), we can prove the Little's theorem as follows:

- ▶ Equate both sides: $\int_0^t N(\tau)d\tau = \sum_{i=1}^{\alpha(t)} T_i$.
- ▶ Divide them by t : $\frac{\int_0^t N(\tau)d\tau}{t} = \frac{\sum_{i=1}^{\alpha(t)} T_i}{t}$.
- ▶ Take the limit of t : $\lim_{t \rightarrow \infty} \frac{\int_0^t N(\tau)d\tau}{t} = \lim_{t \rightarrow \infty} \frac{\sum_{i=1}^{\alpha(t)} T_i}{t}$.
- ▶ Note that $\lim_{t \rightarrow \infty} \frac{\int_0^t N(\tau)d\tau}{t} = N$ and
$$\lim_{t \rightarrow \infty} \frac{\sum_{i=1}^{\alpha(t)} T_i}{t} = \lim_{t \rightarrow \infty} \left[\left(\frac{\sum_{i=1}^{\alpha(t)} T_i}{\alpha(t)} \right) \times \left(\frac{\alpha(t)}{t} \right) \right] = T\lambda.$$
- ▶ Finally, we have $N = \lambda T$.

Little's Theorem: Applications

- ▶ If N_Q is the average number of customers waiting in the queue (but not under service) and W is the average customer waiting time in the queue, then

$$N_Q = \lambda W.$$

- ▶ If \bar{X} is the average service time and ρ is the average number of customers under service (also called “utilization factor” in the communication system), then

$$\rho = \lambda \bar{X}.$$

Next ...

Basic Queueing Theory

Little's Theorem

Poisson Process

M/M/1 Queueing System

Other Markov Systems

Network Simulation

Ergodicity

Time-Stepped vs. Discrete-Event Simulation

Discrete-Event Simulation

SimPy: Python-Based Discrete-Event Simulation

Framework

Poisson Process: Introduction

- ▶ In queueing theory, the arrival process of customers often can be described by a Poisson process.
- ▶ In communication networks, the *customers* may be calls or packets. Poisson process is a viable model when the calls or packets originate from a large population of independent users.
- ▶ Poisson process is mathematically described by a counting process $A(t)$, i.e., for $t > 0$ and $s > t$,
 - ▶ $A(0) = 0$
 - ▶ $A(t)$ = the # of arrivals during the interval $[0, t]$.
 - ▶ $A(t - s) \triangleq A(t) - A(s)$ = the # of arrivals during the interval $(s, t]$.

Poisson Process: Definition 1 - Pure Birth Process

A *Poisson process* with rate λ can be defined in the three different but *equivalent* ways.

- ▶ In an infinitesimal time interval δ , there may occur only one arrival with the probability $\lambda\delta$, i.e., for every $t \geq 0$ and $\delta \geq 0$,

$$P \{A(t + \delta) - A(t) = 0\} = 1 - \lambda\delta + o(\delta),$$

$$P \{A(t + \delta) - A(t) = 1\} = \lambda\delta + o(\delta),$$

$$P \{A(t + \delta) - A(t) \geq 2\} = o(\delta),$$

where $o(\delta)$ is a function such that $\lim_{\delta \rightarrow 0} \frac{o(\delta)}{\delta} = 0$.

- ▶ The arrival is independent of arrivals outside the interval.

Poisson Process: Definition 2 - Process of Independent Increments

- ▶ The number of arrivals in any finite interval of length t follows a Poisson distribution with the average of λt , i.e.,


$$P\{A(t) = n\} = e^{-\lambda t} \frac{(\lambda t)^n}{n!}, \quad n = 0, 1, \dots$$

- ▶ The number of arrivals in disjoint intervals are independent.

Poisson Process: Definition 3 - Exponential Interarrival Time

- Interarrival times are independent and exponentially distributed with parameter λ ; if t_n denotes the time of the n th arrival, the intervals $\tau_n = t_{n+1} - t_n$ have the probability distribution

$$P\{\tau_n \leq s\} = 1 - e^{-\lambda s}, \quad s \geq 0.$$



Poisson Process: Equivalence of Definitions - $1 \rightarrow 2$ I

- From the definition 1, we have

$$P\{A(t + \delta) - A(t)\} = \lambda\delta + o(\delta).$$

- Consider the probability generating function $G_t(z)$:

$$G_t(z) = E\left[z^{A(0,t)}\right],$$

$$\begin{aligned} G_{t+\delta}(z) &= E\left[z^{A(0,t+\delta)}\right] = E\left[z^{A(0,t)+A(t,t+\delta)}\right] \\ &= E\left[z^{A(0,t)}\right] E\left[z^{A(t,t+\delta)}\right] = G_t(z) (1 - \lambda\delta + \lambda\delta z) \\ &= G_t(z) - \lambda\delta(1 - z)G_t(z) \end{aligned}$$

$$\frac{G_{t+\delta}(z) - G_t(z)}{\delta} = \lambda(z - 1)G_t(z).$$

Poisson Process: Equivalence of Definitions - 1 \rightarrow 2 II

- ▶ Taking limit (i.e., $\delta \rightarrow 0$) provides the following differential equation:

$$\begin{aligned}\frac{d}{dt}G_t(z) &= \lambda(z-1)G_t(z), \\ G_t(z) &= e^{(z-1)\lambda t},\end{aligned}$$

where the last equation is the probability generating function of the Poisson distribution.

Poisson Process: Equivalence of Definitions - $2 \rightarrow 3$

- ▶ Note that the the following two events are identical:
 - ▶ An interarrival time is greater than t .
 - ▶ There is no arrival during the time interval of length t .
- ▶ If the interarrival time is for n th and $(n + 1)$ th arrivals (i.e., $\tau_n = t_{n+1} - t_n$), we have

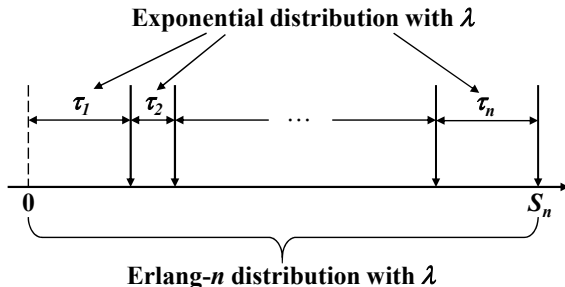
$$P\{\tau_n > t\} = P\{A(t_n + t, t_n) = 0\} = e^{-\lambda t},$$

which shows the interarrival time follows the exponential distribution.

Poisson Process: Equivalence of Definitions - 3 \rightarrow 1 I

- If the interarrival times follow the exponential distribution with parameter λ , the n th arrival time $S_n \triangleq \sum_{i=1}^n \tau_i$ follows Erlang- n distribution¹, i.e.,

$$P\{S_n \leq t\} = 1 - \sum_{i=0}^{n-1} \frac{e^{-\lambda t} (\lambda t)^i}{i!}.$$



Poisson Process: Equivalence of Definitions - 3 \rightarrow 1 II

- ▶ For infinitesimal time δ , therefore, we have the following probabilities:

$$P\{A(t + \delta) - A(t) = 0\} = P\{\tau_n > \delta\} = e^{-\lambda\delta} = 1 - \lambda\delta + o(\delta),$$

$$P\{A(t + \delta) - A(t) \geq 2\} = P\{S_2 \leq \delta\}$$

$$= 1 - \sum_{i=0}^1 \frac{e^{-\lambda\delta}(\lambda\delta)^i}{i!} = o(\delta),$$

where we apply Maclaurin series (i.e., Taylor series centered at zero) expansion of the exponential function.

- ▶ $P\{A(t + \delta) - A(t) = 1\}$ can be derived from the above and the probability normalization condition.

¹<http://www.math.unl.edu/~scohn1/428s05/queue3.pdf>

Poisson Process: Equivalence of Definitions (Optional) - 2 \rightarrow 1

If the counting process $A(t)$ follows the Poisson distribution, i.e., $P\{A(t)=n\} = e^{-\lambda t} \frac{(\lambda t)^n}{n!}$, then

$$P\{A(\delta) = 0\} = e^{-\lambda\delta} = 1 - \lambda\delta + o(\delta),$$

$$P\{A(\delta) = 1\} = \frac{\lambda\delta}{1!} e^{-\lambda\delta} = \lambda\delta + o(\delta).$$

Next ...

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System**

- Other Markov Systems

Network Simulation

- Ergodicity

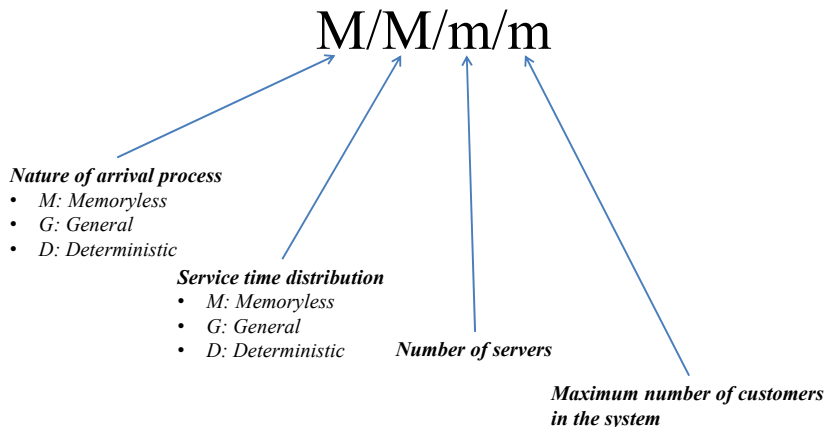
- Time-Stepped vs. Discrete-Event Simulation

- Discrete-Event Simulation


- SimPy: Python-Based Discrete-Event Simulation

- Framework

Queueing Theory Nomenclature



M/M/1 Queueing System: Overview



- Note that successive interarrival and service times are assumed to be *statistically independent of each other*.

M/M/1 Queueing System: Arrival Statistics

- ▶ The customer arrivals follow a Poisson process with rate λ .
 - ▶ The interarrival times, therefore, follow an exponential distribution with parameter λ , i.e.,

$$P\{\tau_n \leq t\} = 1 - e^{-\lambda t}, \quad t \geq 0,$$

where τ_n is the interarrival time between the n th and the $(n + 1)$ th customers.

- ▶ The interarrival times τ_n are mutually independent and also independent of all interarrival times.
- ▶ The parameter λ is called the *arrival rate* and represents the average rate at which the customer arrives at the system.

M/M/1 Queueing System: Service Statistics

- ▶ The customer service times follow an exponential distribution with parameter μ , i.e.,

$$P\{s_n \leq t\} = 1 - e^{-\mu t}, \quad t \geq 0,$$

where s_n is the service time of the n th customer.

- ▶ The service times s_n are mutually independent and also independent of all service times.
- ▶ The parameter μ is called the *service rate* and represents the average rate at which the server operates when busy.

M/M/1 Queueing System: Memoryless Property

- ▶ The exponential distribution has a memoryless property which is expressed as follows (for the interarrival and service times τ_n and s_n , respectively):

$$P\{\tau_n > r + t | \tau_n > t\} = P\{\tau_n > r\}, \text{ for } r, t \geq 0,$$

$$P\{s_n > r + t | s_n > t\} = P\{s_n > r\}, \text{ for } r, t \geq 0.$$

M/M/1 Queueing System: Markov Chain Formulation I

- ▶ The memoryless property allows the use of *Markov chains* (here we use a discrete-time version for simplicity):
 - ▶ We focus only on the number of customers in the system at discrete times $0, \delta, 2\delta, \dots$ (δ is a small positive number).
 - ▶ The set $\{N_k | k=0, 1, \dots\}$ forms a discrete-time Markov chain with transition probabilities $P_{i,j} = P\{N_k=j | N_{k-1}=i\}$, where N_k is defined as $N(k\delta)$.

M/M/1 Queueing System: Markov Chain Formulation II

- Now the transition probabilities are given by


$$P_{0,0} = 1 - \lambda\delta + o(\delta),$$

$$P_{i,i} = 1 - \lambda\delta - \mu\delta + o(\delta), \quad i \geq 1,$$

$$P_{i,i+1} = \lambda\delta + o(\delta), \quad i \geq 0,$$

$$P_{i,i-1} = \mu\delta + o(\delta), \quad i \geq 1,$$

$$P_{i,j} = o(\delta), \quad i \text{ and } j \neq i, i+1, i-1.$$



M/M/1 Queueing System: Stationary Distribution I


- Define steady-state probabilities

$$p_n = \lim_{k \rightarrow \infty} P \{N_k = n\} = \lim_{t \rightarrow \infty} P \{N(t) = n\}.$$

- A local balance (of flows) equation is given by

$$p_{n-1}\lambda\delta + o(\delta) = p_n\mu\delta + o(\delta).$$

- By dividing both sides by δ and taking the limit $\delta \rightarrow 0$, we obtain $p_{n-1}\lambda = p_n\mu$.



M/M/1 Queueing System: Stationary Distribution II

- Now we can obtain $p_n = \rho^n(1-\rho)$ for $n \geq 0$ from the following two equations (if $\rho < 1$):

$$p_n = \rho^n p_0, \quad n = 1, 2, \dots,$$

$$1 = \sum_{n=0}^{\infty} p_n = \sum_{n=0}^{\infty} \rho^n p_0 = \frac{p_0}{1 - \rho},$$

where $\rho = \frac{\lambda}{\mu}$.

M/M/1 Queueing System: Main Results

- ▶ Average number of customers in the system

$$\overline{N} = \frac{\rho}{1 - \rho} = \frac{\lambda}{\mu - \lambda}$$

- ▶ Average delay

$$\overline{T} = \frac{\overline{N}}{\lambda} = \frac{1}{\mu - \lambda}$$


- ▶ Average waiting time in the queue

$$\overline{W} = \frac{1}{\mu - \lambda} - \frac{1}{\mu} = \frac{\rho}{\mu - \lambda}$$

- ▶ Average number of customers in the queue

$$\overline{N}_Q = \lambda \overline{W} = \frac{\rho^2}{1 - \rho}$$

M/M/1 Queueing System: Average Delay



Next ...

Basic Queueing Theory

Little's Theorem

Poisson Process

M/M/1 Queueing System

Other Markov Systems

Network Simulation


Ergodicity

Time-Stepped vs. Discrete-Event Simulation


Discrete-Event Simulation

SimPy: Python-Based Discrete-Event Simulation
Framework

Other Markov Systems



M/M/m – The System with m Servers



M/M/m/m – The m -Server Loss System

M/M/m Queueing System: Stationary Distribution I

- From the local balance equations, we get the following (if $\rho \triangleq \lambda/(m\mu) < 1$):

$$\begin{cases} p_n = \frac{\lambda}{n\mu} p_{n-1}, & \text{for } 1 \leq n < m, \\ p_n = \frac{\lambda}{m\mu} p_{n-1}, & \text{for } n \geq m. \end{cases}$$

- Now we can obtain p_n as follows:

$$\begin{cases} p_n = \frac{(m\rho)^n}{m^n} p_0, & \text{for } 0 \leq n < m, \\ p_n = \frac{m^n \rho^n}{m!} p_0, & \text{for } n \geq m. \end{cases}$$

M/M/m Queueing System: Stationary Distribution II

- ▶ Because $\sum_{n=0}^{\infty} p_n = 1$, we obtain

$$p_0 = \frac{1}{\sum_{n=0}^{m-1} \frac{(m\rho)^n}{n!} + \frac{(m\rho)^m}{m!} \frac{1}{1-\rho}}$$

- ▶ We first obtain p_{m+} , which will be used in deriving main results:

$$p_{m+} \triangleq P\{N \geq m\} = \sum_{n=m}^{\infty} p_n = \frac{(m\rho)^m}{m!(1-\rho)} p_0$$

- ▶ p_{m+} is the probability that all servers are busy and the customers has to wait in the queue (also called *Erlang C* formula).

M/M/m Queueing System: Main Results

- ▶ Average number of customers in the system

$$\bar{N} = m\rho + \frac{\rho}{1-\rho}p_{m+}$$

- ▶ Average number of customers in the queue

$$\bar{N}_Q = \frac{\rho}{1-\rho}p_{m+}$$

- ▶ Average waiting time in the queue

$$\bar{W} = \frac{\bar{N}_Q}{\lambda} = \frac{\rho}{\lambda(1-\rho)}p_{m+}$$

- ▶ Average delay

$$\bar{T} = \bar{W} + \frac{1}{\mu} = \frac{\rho}{\lambda(1-\rho)}p_{m+} + \frac{1}{\mu}$$

Next ...

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System

- Other Markov Systems

Network Simulation

- Ergodicity

- Time-Stepped vs. Discrete-Event Simulation

- Discrete-Event Simulation

- SimPy: Python-Based Discrete-Event Simulation Framework

Network Simulation

Next ...

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System

- Other Markov Systems

Network Simulation

- Ergodicity

- Time-Stepped vs. Discrete-Event Simulation

- Discrete-Event Simulation

- SimPy: Python-Based Discrete-Event Simulation

- Framework

Ergodicity: Introduction

In econometrics and signal processing, a stochastic process is said to be *ergodic* if its statistical properties can be deduced from a single, sufficiently long, random sample of the process.²

- ▶ The reasoning is that any collection of random samples from a process must represent the average statistical properties of the entire process.
- ▶ In other words, regardless of what the individual samples are, a birds-eye view of the collection of samples must represent the whole process.
- ▶ Conversely, a process that is not ergodic is a process that changes erratically at an inconsistent rate.

²The contents on ergodicity are from [Wikipedia](#).

Ergodicity: Mean-Ergodic Process

The process $X(t)$ is said to be *mean-ergodic* or *mean-square ergodic in the first moment* if the time average estimate

$$\hat{\mu}_X = \frac{1}{T} \int_0^T X(t) dt$$

converges in squared mean to the ensemble average μ_X as $T \rightarrow \infty$.

Ergodicity: Autocovariance-Ergodic Process

Likewise, the process is said to be *autocovariance-ergodic* or *mean-square ergodic in the second moment* if the time average estimate

$$\hat{r}_X(\tau) = \frac{1}{T} \int_0^T [X(t + \tau) - \mu_X][X(t) - \mu_X] dt$$

converges in squared mean to the ensemble average $r_X(\tau)$, as $T \rightarrow \infty$.

- ▶ A process which is ergodic in the mean and autocovariance is sometimes called *ergodic in the wide sense*.

Next ...

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System

- Other Markov Systems

Network Simulation

- Ergodicity

- Time-Stepped vs. Discrete-Event Simulation**

- Discrete-Event Simulation


- SimPy: Python-Based Discrete-Event Simulation

- Framework

Time-Stepped vs. Discrete-Event Simulation


Time-Stepped

- ▶ System states are globally updated at periodic time instances with increment dt .
- ▶ Entities exchange state updates via *messages*.



Discrete-Event

- ▶ System states are locally and asynchronously updated at different times.
- ▶ Entities exchange *events* for state updates.



Next ...

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System

- Other Markov Systems

Network Simulation

- Ergodicity

- Time-Stepped vs. Discrete-Event Simulation

- Discrete-Event Simulation**

- SimPy: Python-Based Discrete-Event Simulation Framework

Discrete-Event Simulation: Components

- ▶ **Event list:** List of future events sorted by their event times
- ▶ **Clock:** Unlike real-time simulation, time hops from one event to another.
- ▶ **Random-number generators:** For generating random variables of various kinds based on PRNGs.
- ▶ **Statistics:** Mean waiting time, queue length, . . .
- ▶ **Ending condition:** Run time (e.g., 3 hours) or the number events (e.g., after 10,000 packets generated).

Discrete-Event Simulation: Program Structure

- ▶ Start
 - ▶ Set *Ending Condition* to FALSE.
 - ▶ Initialize clock (typically set to zero) and variables for system states and statistics.
 - ▶ Schedule *the first event* to trigger the simulation.
- ▶ While *Ending Condition* is **FALSE**, do the following:
 - ▶ Fetch *the next event* from the event list.
 - ▶ Set the clock to the time of the fetched event.
 - ▶ Process the event.
 - ▶ Update variables for system states and statistics.
 - ▶ Schedule/cancel future events.
- ▶ End
 - ▶ Generate statistical report

Next ...

Basic Queueing Theory

- Little's Theorem

- Poisson Process

- M/M/1 Queueing System

- Other Markov Systems

Network Simulation

- Ergodicity

- Time-Stepped vs. Discrete-Event Simulation

- Discrete-Event Simulation

- SimPy: Python-Based Discrete-Event Simulation Framework**

SimPy: Overview

SimPy³ is a process-based discrete-event simulation framework based on standard Python⁴.

- ▶ *Processes* are defined by Python *generator functions*⁵ and may, for example, be used to model active components like packets and customers.
- ▶ *Shared resources* are also provided to model limited capacity congestion points like servers and checkout counters.

³<https://simpy.readthedocs.io/>

⁴<https://www.python.org/>

⁵<https://wiki.python.org/moin/Generators>

SimPy: An Illustrating Example

Simulates two clocks ticking in different time intervals.

```
1 >>> import simpy
2 >>>
3 >>> def clock(env, name, tick):
4 ...     while True:
5 ...         print(name, env.now)
6 ...         yield env.timeout(tick)
7 ...
8 >>> env = simpy.Environment()
9 >>> env.process(clock(env, 'fast', 0.5))
10 <Process(clock) object at 0x...>
11 >>> env.process(clock(env, 'slow', 1))
12 <Process(clock) object at 0x...>
13 >>> env.run(until=2)
14 fast 0
15 slow 0
16 fast 0.5
17 slow 1
18 fast 1.0
19 fast 1.5
```

Generator I

```
1 >>> itr = [x*x for x in range(10)]
2 >>> itr
3 [0, 1, 4, 9, 16, 25, 36, 49, 64, 81]
4 >>> len(itr)
5 10
6 >>> gen = (x*x for x in range(10))
7 >>> gen
8 <generator object <genexpr> at 0x0000001532A8B9BF8>
9 >>> len(gen)
10 Traceback (most recent call last):
11   File "<stdin>", line 1, in <module>
12   TypeError: object of type 'generator' has no len()
```

The generator expression for `gen` above is equivalent to the following generator implementation:

```
1 def squared_series():
2     for x in range(10):
3         yield x*x
```

Generator II

Can we make an infinite sequence? Yes, sort of ...

```
1 def pi_series():
2     sum = 0
3     i = 1.0; j = 1
4     while(1):
5         sum = sum + j/i
6         yield 4*sum
7         i = i + 2; j = j * -1
```

This generator generates a series converges to $\pi/4$.

Generator III

How can we use a generator?

- By next() call.


```
1 >>> gen = (x*x for x in range(10))
2 >>> next(gen)
3 0
4 >>> next(gen)
5 1
6 >>> next(gen)
7 4
```

- Inside a loop.

```
1 s = squared_series()
2 for i in s:
3     print(i)
```

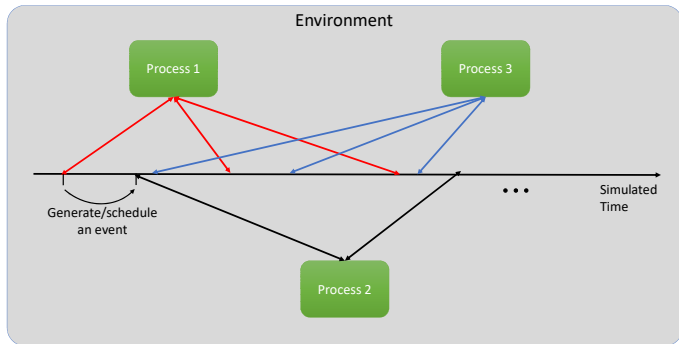
Generator IV

The following figure illustrates the execution of a generator function during its 1st call and calls thereafter.



SimPy: Basic Concepts I

- ▶ The behavior of an active component (e.g., a packet) is modeled with processes. All processes live in an environment. They interact with the environment and with each other via events.



SimPy: Basic Concepts II

- ▶ Processes are described by simple Python generators. During their lifetime, they create events and yield them in order to wait for them to be triggered.
- ▶ When a process yields an event, the process gets suspended. SimPy resumes the process, when the event occurs (we say that the event is triggered).
- ▶ An important event type is the `Timeout`. Events of this type are triggered after a certain amount of (simulated) time has passed. They allow a process to sleep (or hold its state) for the given time. A `Timeout` and all other events can be created by calling the appropriate method of the `Environment` that the process lives in (`Environment.timeout()` for example).

SimPy: Process Interaction

The `Process` instance that is returned by `Environment.process()` can be utilized for process interactions, including

- ▶ Waiting for another process to terminate
- ▶ Interrupting another process

SimPy: Shared Resources

SimPy offers three types of **resources** that help you modeling problems, where multiple processes want to use a resource of limited capacity (e.g., packets at a queueing system with a limited number of servers) or classical producer-consumer problems.

- ▶ *Resources* can be used by a limited number of processes at a time (e.g., a gas station with a limited number of fuel pumps).
- ▶ *Containers* model the production and consumption of a homogeneous, undifferentiated bulk. It may either be continuous (like water) or discrete (like apples).
- ▶ *Stores* allow the production and consumption of Python objects.

Example: M/M/1 Queue I

```
1  #!/usr/bin/env python
2  # -*- coding: utf-8 -*-
3  ##
4  # @file      mml.py
5  # @author    Kyeong Soo (Joseph) Kim <kyeongsoo.kim@gmail.com>
6  # @date      2016-09-27
7  #
8  # @brief     Simulate M/M/1 queueing system
9  #
10 # @remarks   Copyright (C) 2016 Kyeong Soo (Joseph) Kim. All rights reserved.
11 #
12 # @remarks   This software is written and distributed under the GNU General
13 #             Public License Version 2 (http://www.gnu.org/licenses/gpl-2.0.html).
14 #             You must not remove this notice, or any other, from this software.
15 #
16
17 import argparse
18 import numpy as np
19 import random
20 import simpy
21 import sys
22
23
24 def source(env, mean_ia_time, mean_srv_time, server, wait_times, number, trace):
25     """Generates packets with exponential interarrival time."""
26     for i in range(number):
27         ia_time = random.expovariate(1.0 / mean_ia_time)
28         srv_time = random.expovariate(1.0 / mean_srv_time)
29         pkt = packet(env, 'Packet-%d' % i, server, srv_time, wait_times, trace)
30         env.process(pkt)
31         yield env.timeout(ia_time)
```

Example: M/M/1 Queue II

```
32
33
34 def packet(env, name, server, service_time, wait_times, trace):
35     """Requests a server, is served for a given service_time, and leaves the server."""
36     arrv_time = env.now
37     if trace:
38         print('t=%.4Es: %s arrived' % (arrv_time, name))
39
40     with server.request() as request:
41         yield request
42
43         wait_time = env.now - arrv_time
44         wait_times.append(wait_time)
45         if trace:
46             print('t=%.4Es: %s waited for %.4Es' % (env.now, name, wait_time))
47         yield env.timeout(service_time)
48         if trace:
49             print('t=%.4Es: %s served for %.4Es' %
50                 (env.now, name, service_time))
51
52
53 def run_simulation(mean_ia_time, mean_srv_time, num_packets=1000, random_seed=1234, trace=True):
54     """Runs a simulation and returns statistics."""
55     print('M/M/1 queue\n')
56     random.seed(random_seed)
57     env = simpy.Environment()
58
59     # start processes and run
60     server = simpy.Resource(env, capacity=1)
61     wait_times = []
```


Example: M/M/1 Queue III

```
62     env.process(source(env, mean_ia_time,
63                       mean_srv_time, server, wait_times, number=num_packets, trace=trace))
64     env.run()
65
66     # return statistics (i.e., mean waiting time)
67     return np.mean(wait_times)
68
69
70 if __name__ == "__main__":
71     parser = argparse.ArgumentParser()
72     parser.add_argument(
73         "-A",
74         "--mean_ia_time",
75         help="mean packet interarrival time [s]; default is 1.0",
76         default=1.0,
77         type=float)
78     parser.add_argument(
79         "-S",
80         "--mean_srv_time",
81         help="mean packet service time [s]; default is 0.1",
82         default=0.1,
83         type=float)
84     parser.add_argument(
85         "-N",
86         "--num_packets",
87         help="number of packets to generate; default is 1000",
88         default=1000,
89         type=int)
90     parser.add_argument(
91         "-R",
```

Example: M/M/1 Queue IV

```
92     "--random_seed",
93     help="seed for random number generation; default is 1234",
94     default=1234,
95     type=int)
96 parser.add_argument('--trace', dest='trace', action='store_true')
97 parser.add_argument('--no-trace', dest='trace', action='store_false')
98 parser.set_defaults(trace=True)
99 args = parser.parse_args()
100
101 # set variables using command-line arguments
102 mean_ia_time = args.mean_ia_time
103 mean_srv_time = args.mean_srv_time
104 num_packets = args.num_packets
105 random_seed = args.random_seed
106 trace = args.trace
107
108 # run a simulation
109 mean_waiting_time = run_simulation(
110     mean_ia_time, mean_srv_time, num_packets, random_seed, trace)
111
112 # print statistics from the simulation
113 print("Average waiting time = %.4Es\n" % mean_waiting_time)
```

Development Environment

- ▶ Download and install **WinPython** distribution.
- ▶ Download a sample program **"mm1.py"** from the ICE.
- ▶ Open "WinPython Command Prompt" and type the following command:
 - ▶ **"python mm1.py"**
- ▶ Congratulations! You just installed SimPy and run your first simulation program.
- ▶ For a better development and debugging, we would recommend the following:
 - ▶ **ipython** within the "WinPython Command Prompt"
 - ▶ **Spyder**, a graphical IDE

References I

- [1] L. Kleinrock, *Queueing Systems - Volume I; Theory*.
John Wiley & Sons, 1974.